

Text Encoding Methods in NLP

1. **Bag of Words (BoW)**

Concept: Represents text by counting the frequency of each word in a document. It treats each document as an unordered collection or "bag" of words.

- **Advantages**: Simple to understand and implement; suitable for systems where word order is less relevant.
- **Disadvantages**: Loses all information about word order (syntax) and context, resulting in a failure to capture meanings dependent on the sequence of words.

Example: Consider the text "dog bites man" and "man bites dog."

- Vocabulary: {dog, bites, man}
- BoW Vectors:
 - "dog bites man": [1, 1, 1]
 - "man bites dog": [1, 1, 1]
- Despite the different meanings, both sentences have the same representation.

2. **N-gram**

Concept: Extends BoW by considering sequences of 'n' consecutive words, thus capturing some local context.

- **Advantages**: Better at capturing phrases and common collocations, slightly more sensitive to the order of words.
- **Disadvantages**: Still results in high dimensionality and sparsity, especially for large 'n'. It increases the computational complexity.

Example: Using the bigram model on the text "the quick brown fox."

- Vocabulary for Bigrams: {the quick, quick brown, brown fox}
- Bigram Vector: [1, 1, 1]
- This provides a peek into the word order but still misses wider sentence context.

3. **One-Hot Encoding**

Concept: Each word is represented as a binary vector where only one element is '1', and all others are '0'. The position of '1' corresponds to the word's position in the vocabulary.

- **Advantages**: Unambiguous representation of words.

- **Disadvantages**: Very high dimensional space with sparse vectors; no similarity between words is captured (e.g., "king" and "monarch" are as different as "king" and "book").

Example: For words "king" and "queen" in a vocabulary {king, queen, royal}

- One-hot for "king": [1, 0, 0]

- One-hot for "queen": [0, 1, 0]

4. **TF-IDF (Term Frequency-Inverse Document Frequency)**

Concept: Weighs the word counts by a measure of how common they are in the corpus to help identify the most important words in a document.

- **Advantages**: Reduces the weight of terms that occur very frequently, thus giving more importance to terms that are more specific to a particular document.

- **Disadvantages**: Still relies on simple term frequencies; ignores word order and context.

Example: In a corpus where "dog" is very common but "bites" is rare, TF-IDF would assign a higher weight to "bites."

5. **Word2Vec**

Concept: Uses neural networks to learn word associations from a large corpus of text. Typically configured in one of two ways: Skip-gram (predicts surrounding words from a current word) or CBOW (predicts a word given its context).

- **Advantages**: Captures a rich amount of word relationships, such as synonyms, antonyms, and more complex linguistic patterns.

- **Disadvantages**: Does not account for word meanings changing based on context (polysemy).

Example: "Paris" and "France" will have vectors that place them closer together than "Paris" and "apple," reflecting their semantic relationship.

6. **GloVe (Global Vectors)**

Concept: Constructs a global word-word co-occurrence matrix from a corpus and then derives the lower-dimensional word vectors.

- **Advantages**: Incorporates both global statistics of the corpus and local statistics of individual documents, capturing both macro and micro-level word relationships.

- **Disadvantages**: Like Word2Vec, does not handle polysemy well.

Example: Similar to Word2Vec, words with similar meanings are embedded close to each other.

7. **FastText**

Concept: An extension of Word2Vec that represents words as bags of character n-grams, which allows it to capture the morphology of words.

- **Advantages**: Can generate vectors for out-of-vocabulary words by summing up the n-gram vectors, thus solving a major limitation of Word2Vec and GloVe.

- **Disadvantages**: Slower to train than Word2Vec due to increased complexity from handling n-grams.

Example: For "apple", character n-grams might include: "ap", "pp", "pl", "le", "apple". If "apple" wasn't seen during training but its n-grams were, FastText can still provide a reasonable vector.

8. **BERT (Bidirectional Encoder Representations from Transformers)**

Concept: Utilizes the Transformer architecture to model language bidirectionally, considering all surrounding text to

generate word embeddings dynamically based on context.

- **Advantages**: Effective at understanding context and nuances of language, providing state-of-the-art results in many NLP tasks.

- **Disadvantages**: Computationally expensive to train and fine-tune; slower inference compared to non-contextual models.

Example: In "He went to the bank to get money" and "He sat by the river bank", "bank" would have different vectors in each context.

9. **GPT (Generative Pre-trained Transformer)**

Concept: An autoregressive language model that generates text based on the probability of occurrence of a word given all previous words.

- **Advantages**: Capable of generating coherent and contextually relevant text across various domains without needing task-specific data.

- **Disadvantages**: Resource-intensive to train; can generate nonsensical or biased text.

Example: Given "The climate of Mars", GPT might generate ", although very cold and harsh, has been a subject of scientific study due to its potential to support human life in the future."

10. **Transformer-XL**

Concept: An improvement over traditional Transformers, capable of handling much longer sequences of data by retaining information from previous segments of text.

- **Advantages**: Allows for more effective modeling of long-term dependencies in text, useful in applications like document summarization or script generation.

- **Disadvantages**: More complex and computationally intensive than standard Transformers.

Example: In a long document, Transformer-XL can remember and use context from several paragraphs back, critical for understanding subsequent sections.

11. **ELMo (Embeddings from Language Models)**

Concept: Uses a deep, bidirectional LSTM network to create dynamic word embeddings based on the context across the entire sentence.

- **Advantages:** Produces context-dependent embeddings that significantly improve performance on many NLP tasks.
- **Disadvantages:** Less efficient computationally compared to Transformer-based models due to the sequential nature of LSTMs.

Example: "stick" would have different embeddings in "a stick of butter" and "stick the poster on the wall."